

Evolving systems design

from unreliable rpc to resilience with **Linkerd**

Edward Wilde

3 April, 2018

FORM³



FINANCIAL CLOUD



G. Pascal
Zachary

SHOW- STOPPER!

The
Breakneck
Race
to Create
Windows NT
and the
Next
Generation
at Microsoft

A story about unreliable RPC

- System 1

- Increased reliability
- Reduced overall latency

But....

- Did not manage to reduce tail-latency
- Did not manage to protect services

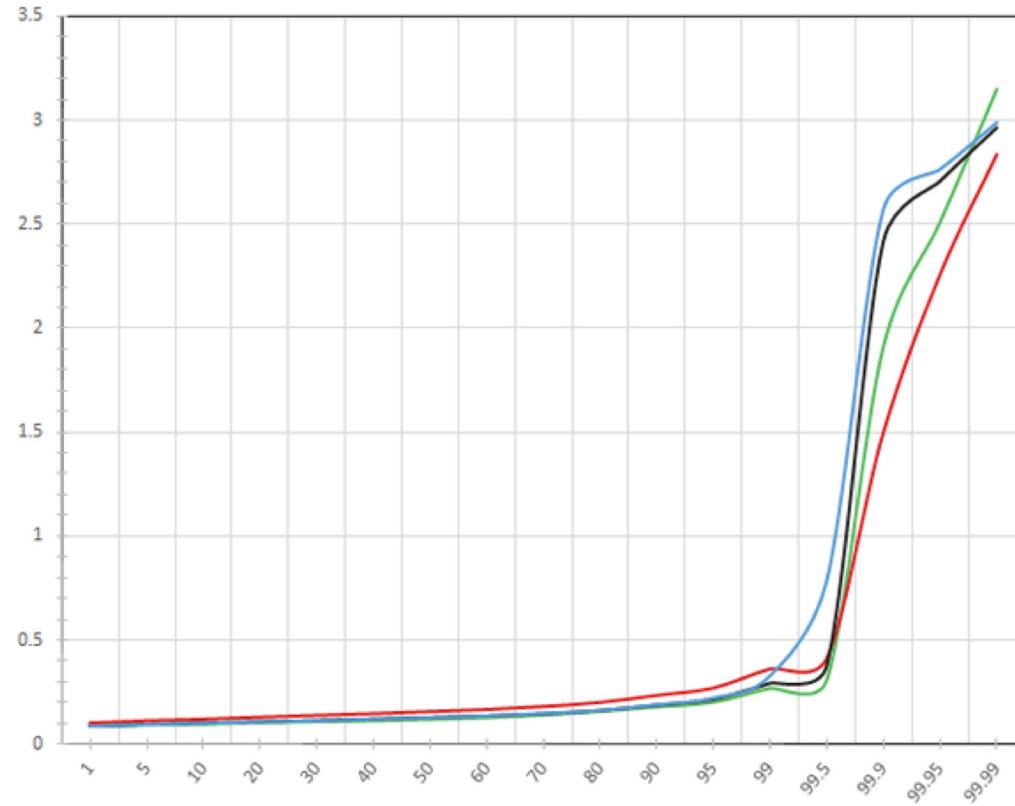
- System 2

- Increased reliability
- Reduced overall latency
- Protected services

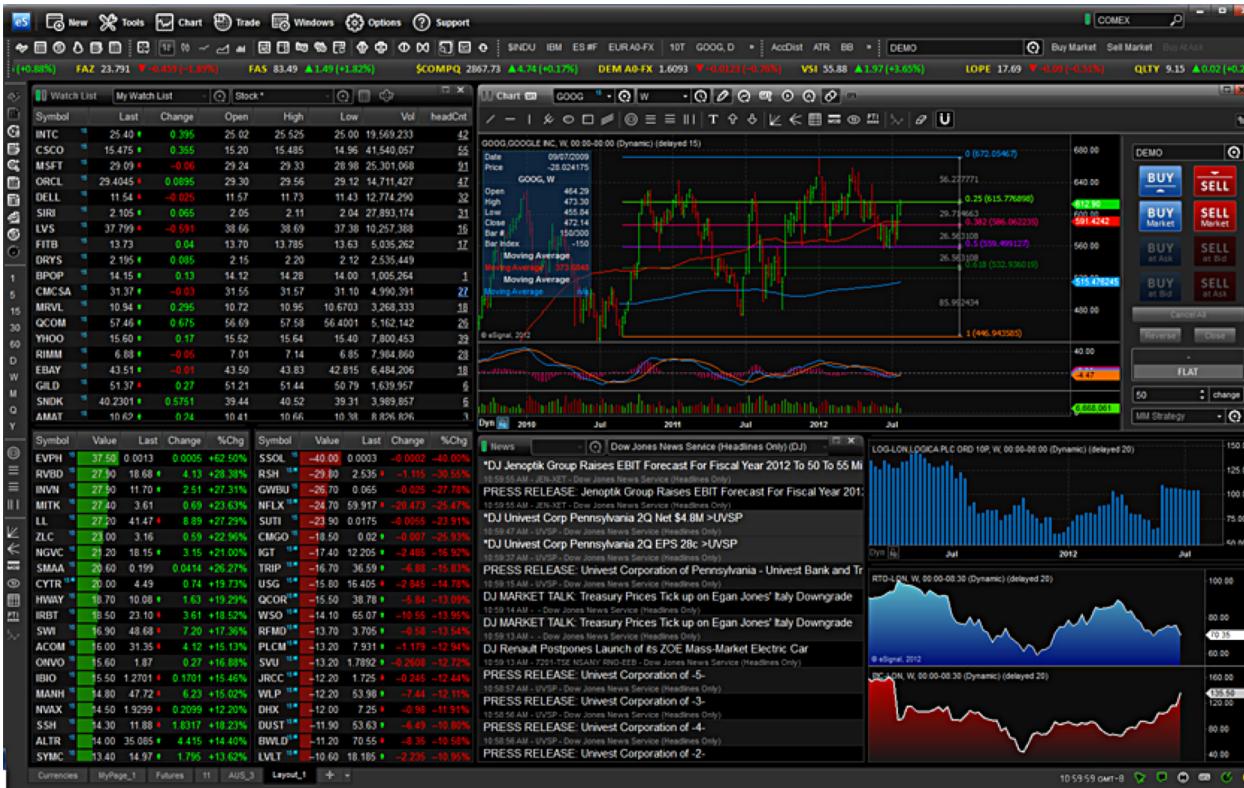
But...

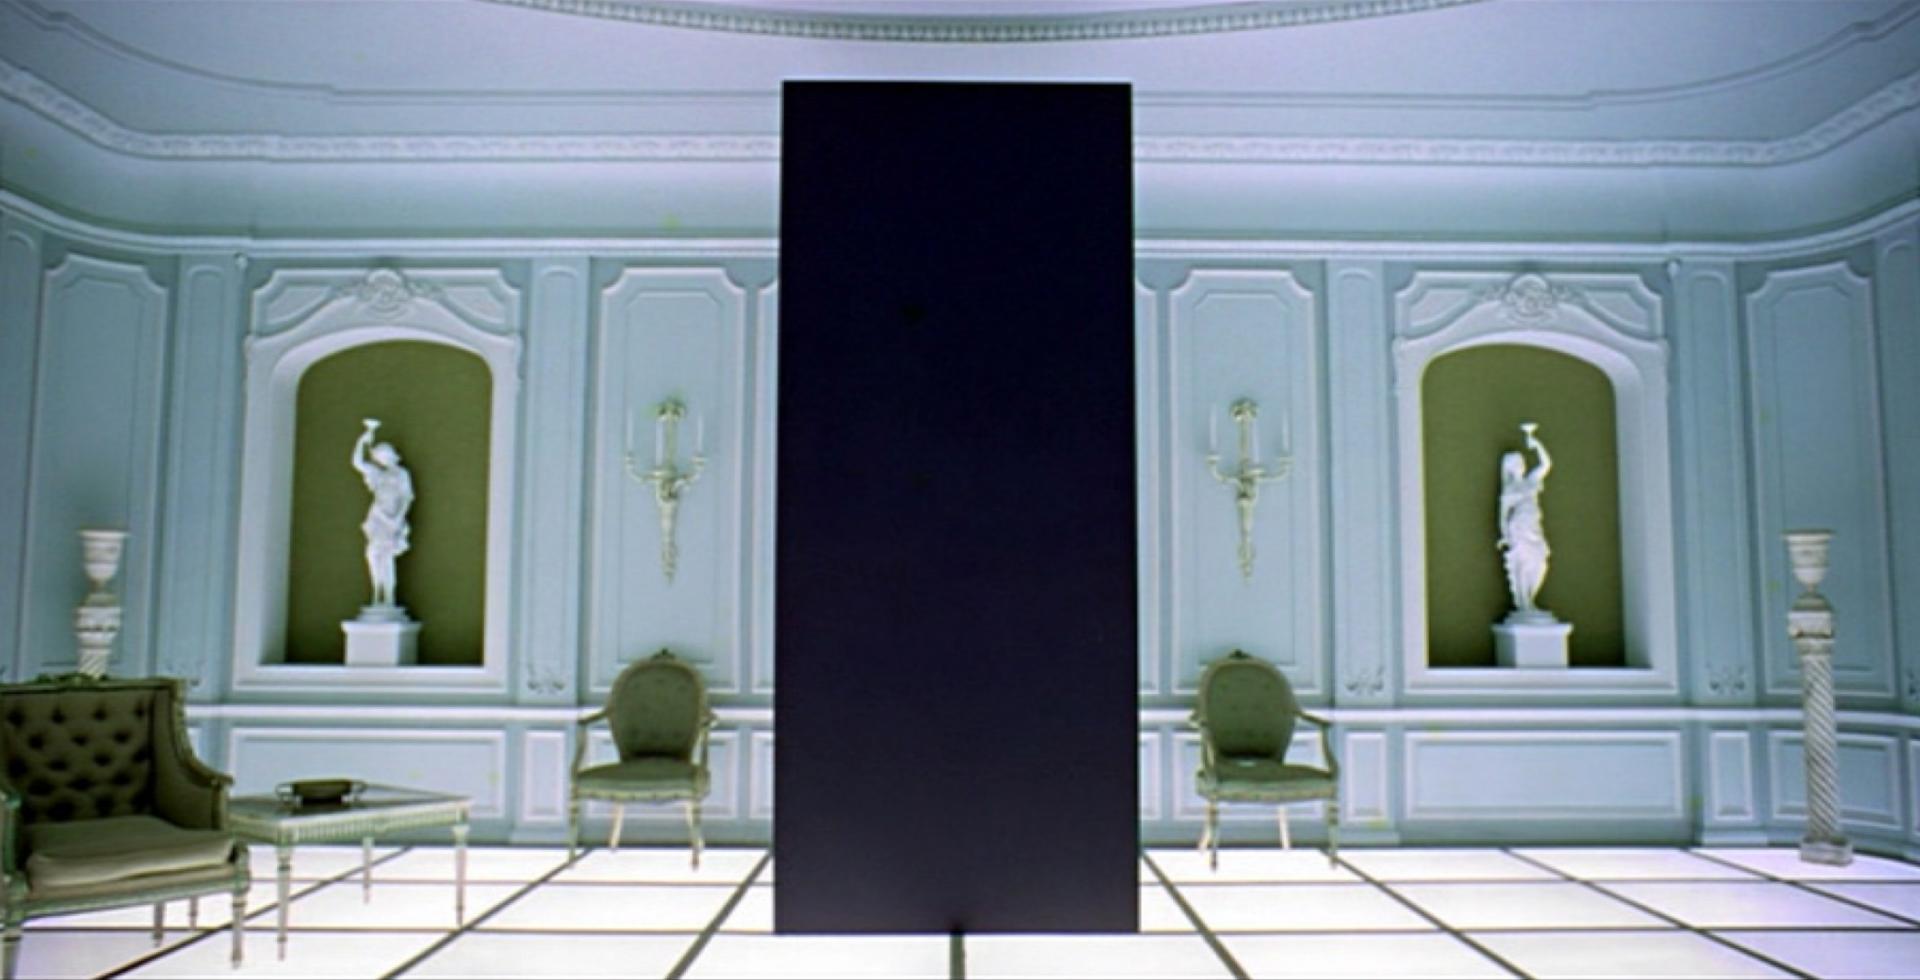
- Still need to work more on tail-latency

Tail latency

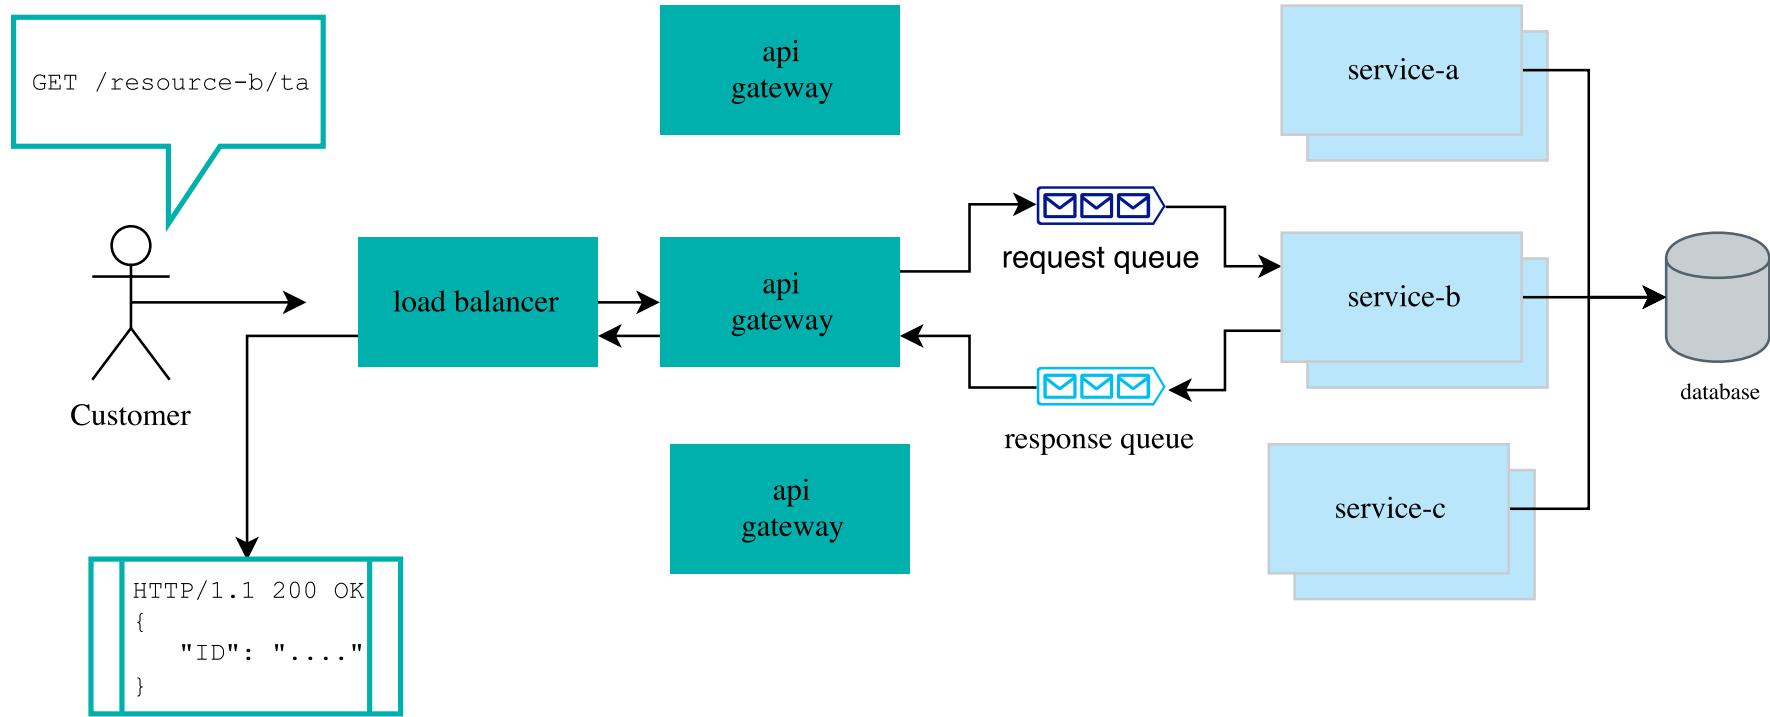


System 1

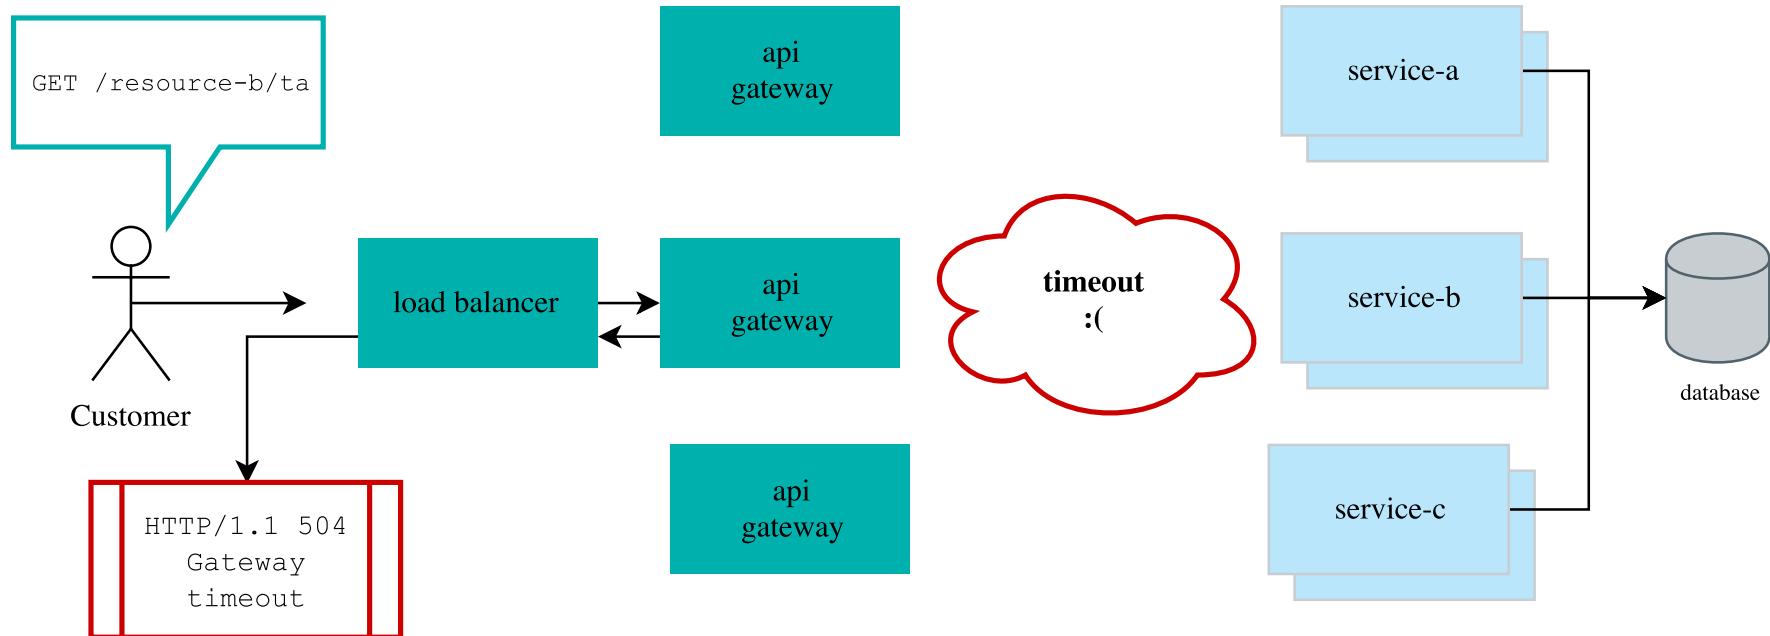




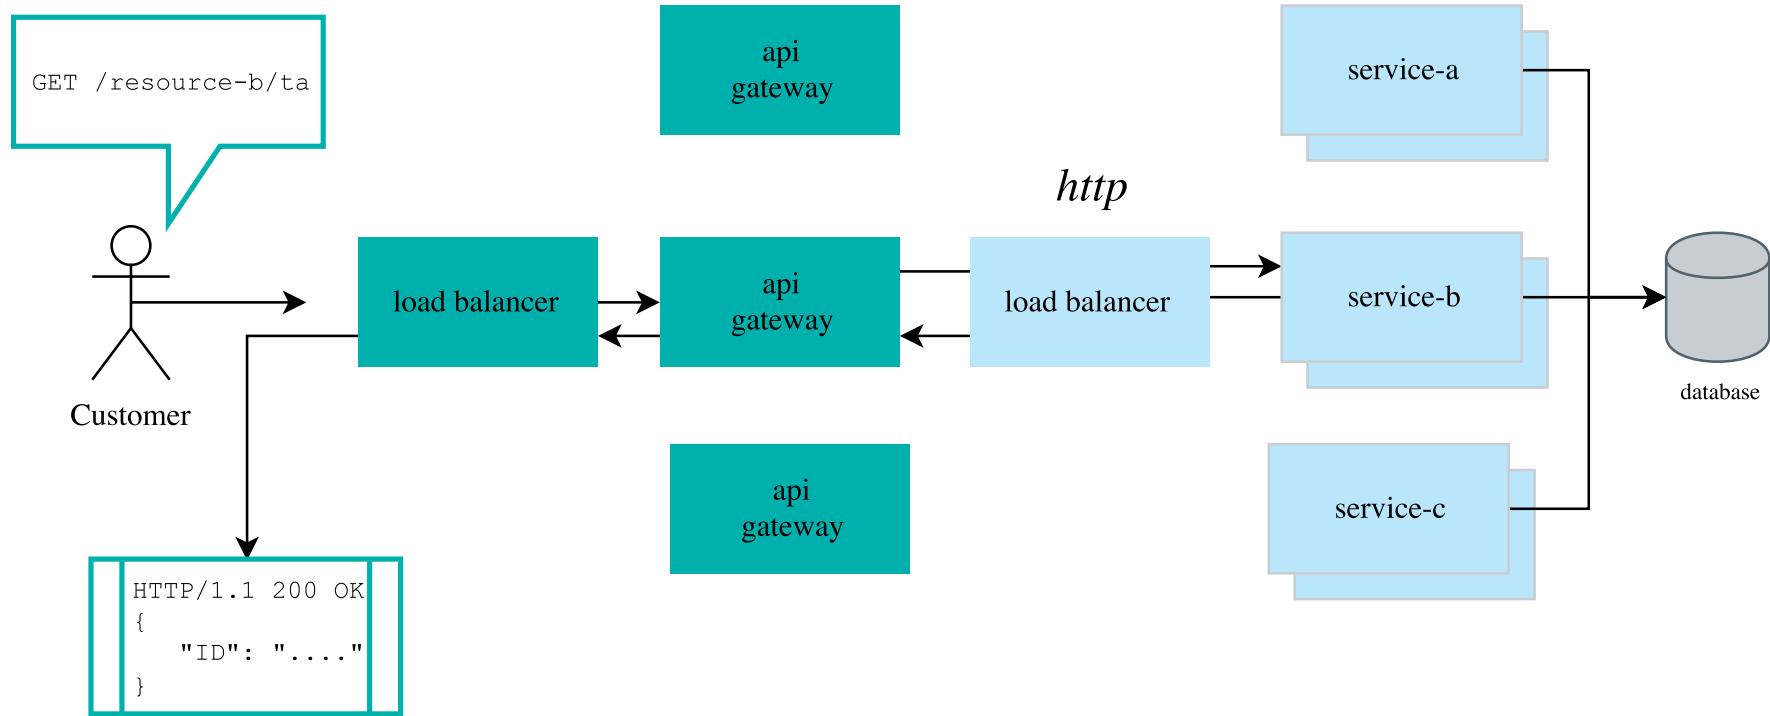
System 1



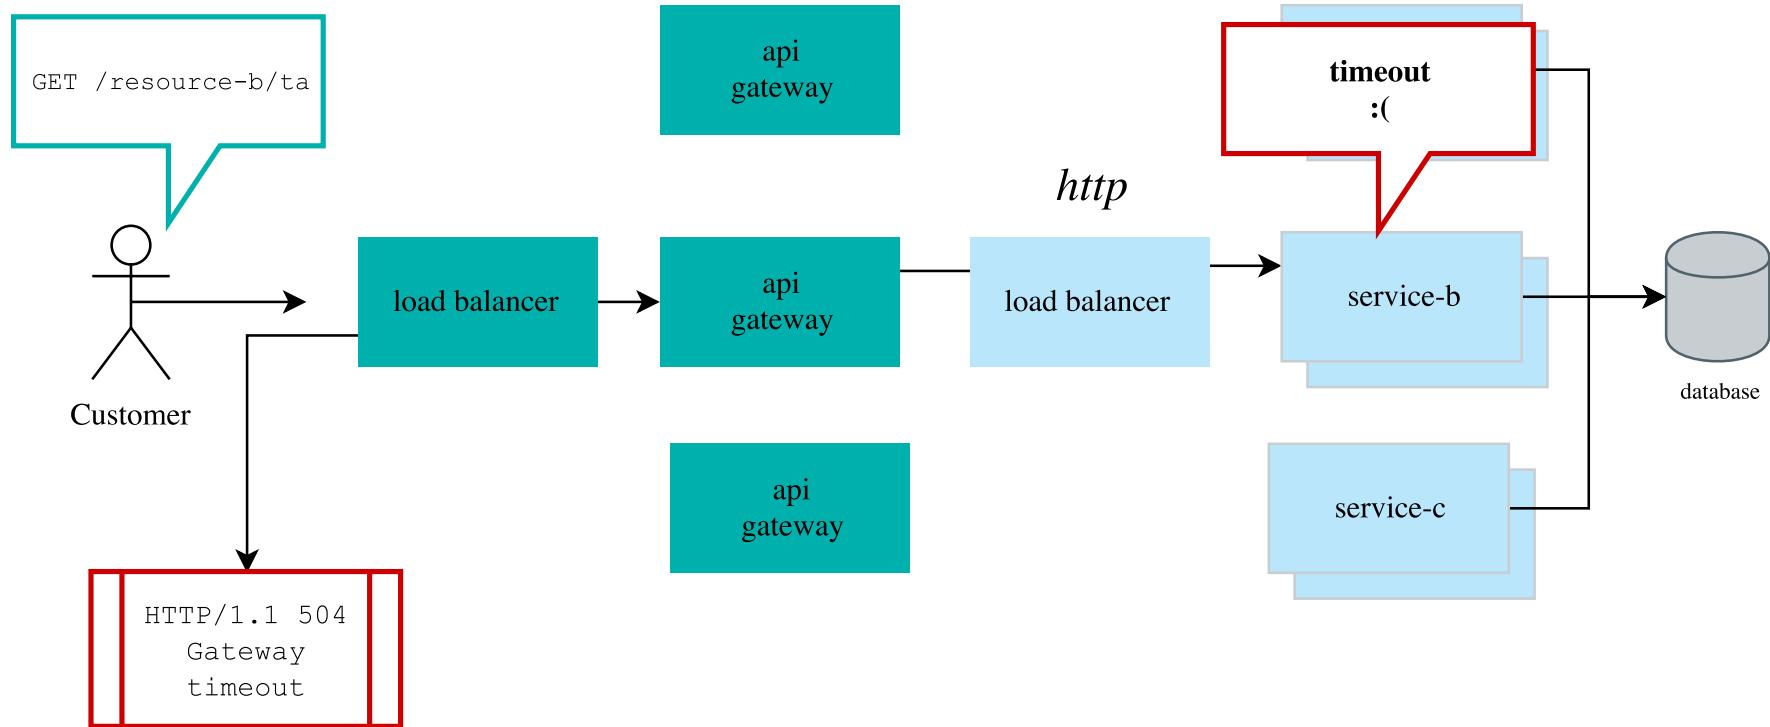
Timeout!



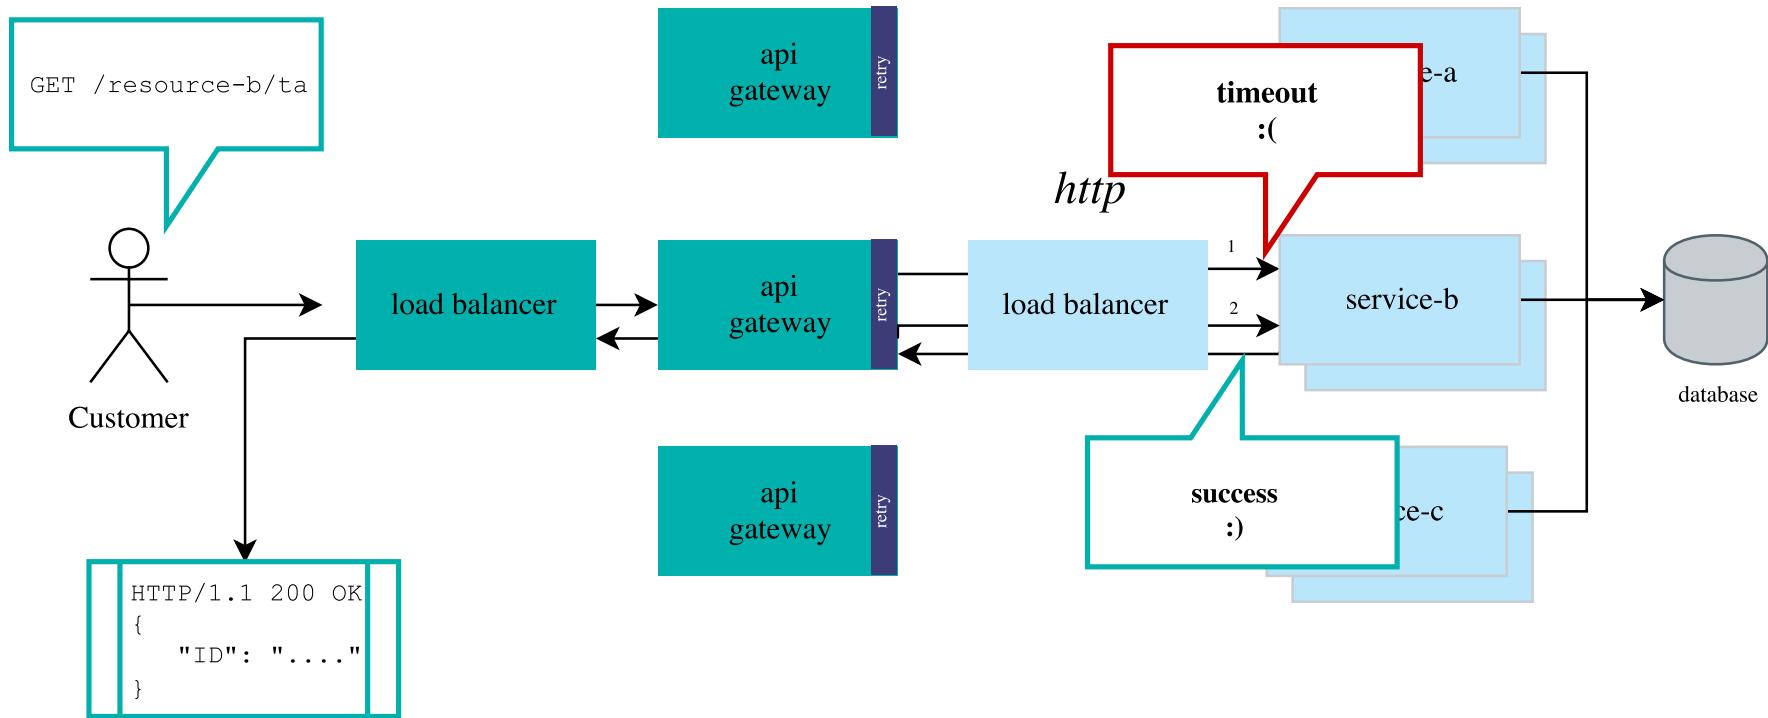
Let's try http



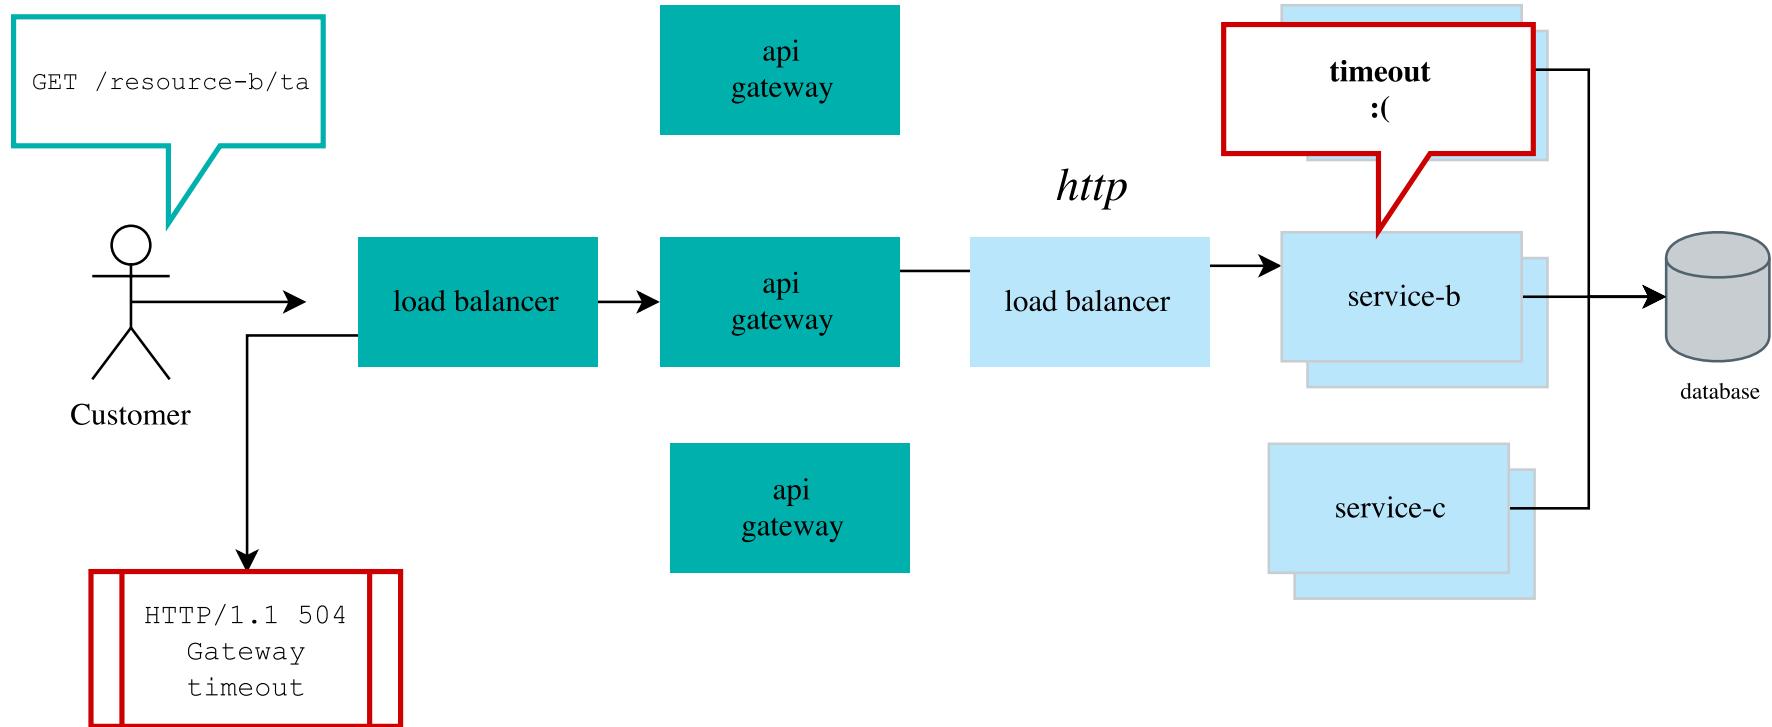
Oops still timeout



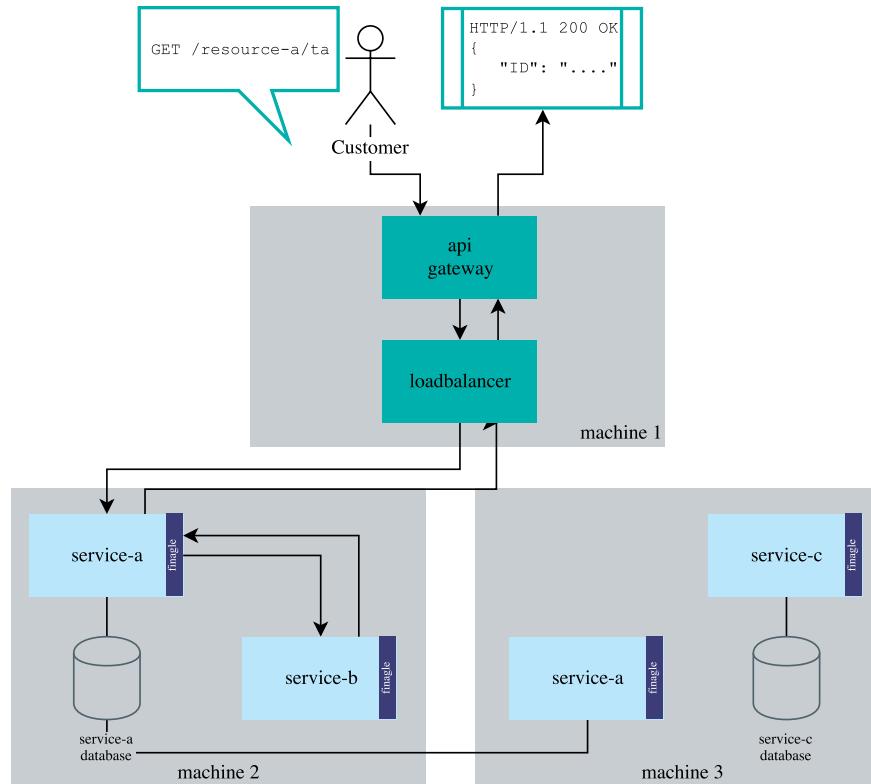
Try that again



Busy day, more timeout



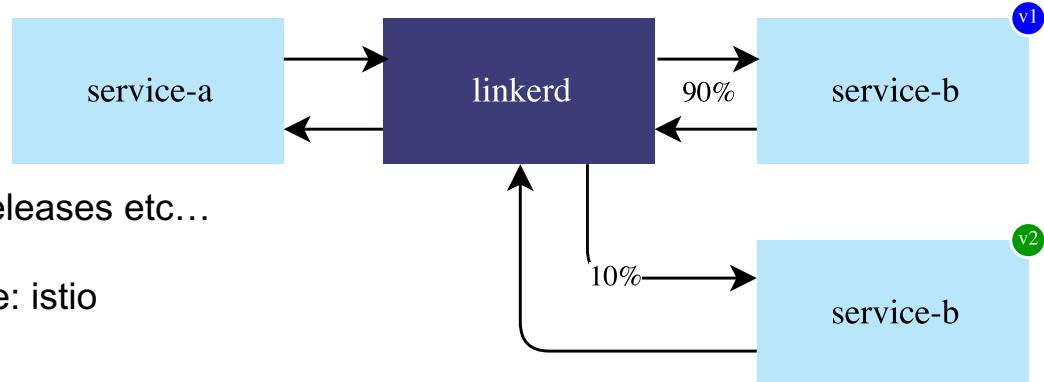
System 2: Service to service communication in a microservices architecture



Service mesh



- Retry failures / timeout
- Implements circuit breaker
- Security
- Routing logic & service discovery
- Blue / green deployments: canary releases etc...
- Distributed Tracing
- Other services meshes are available: istio



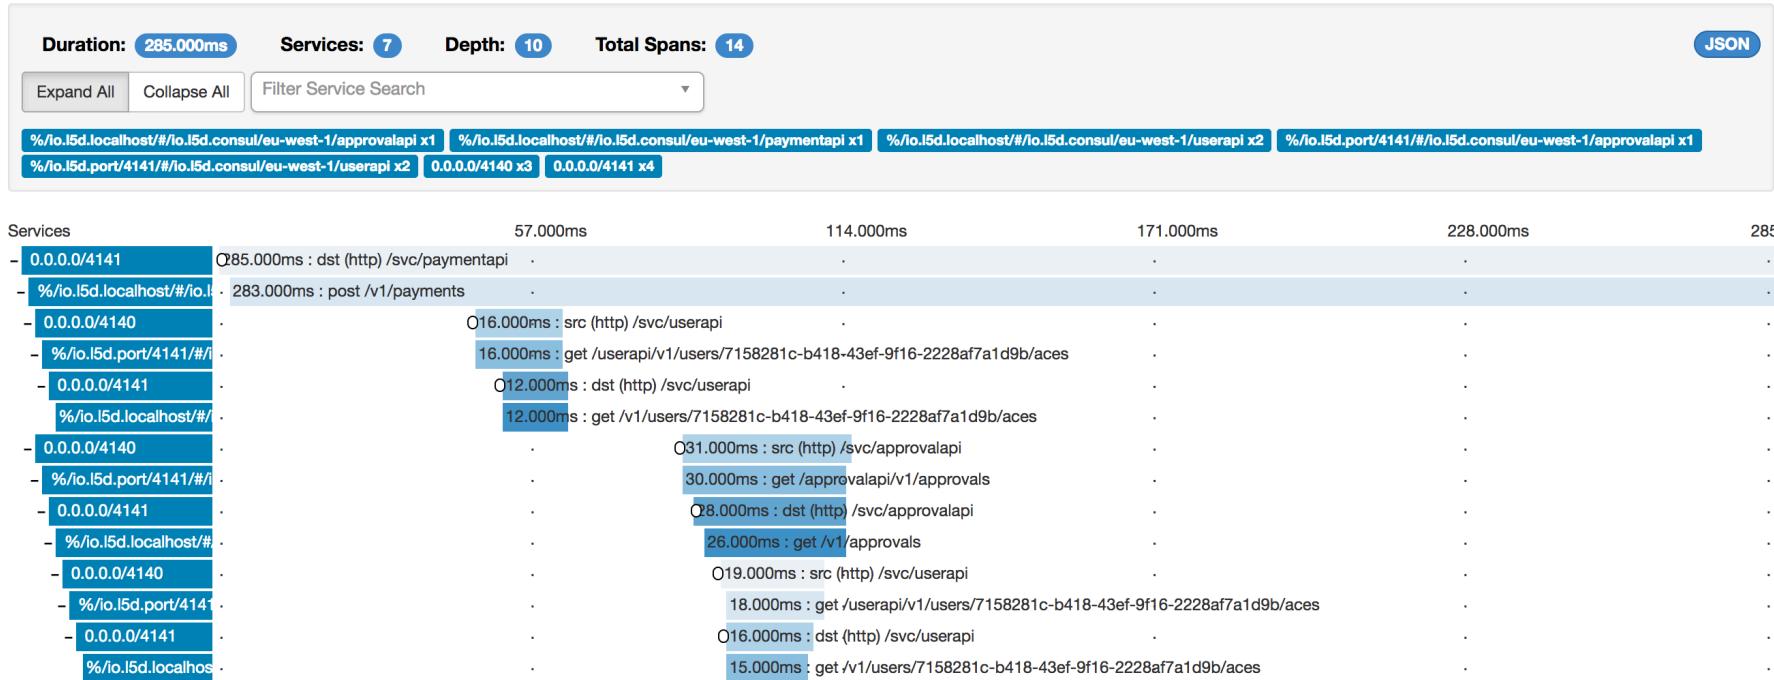
Tracing

Zipkin Investigate system behavior

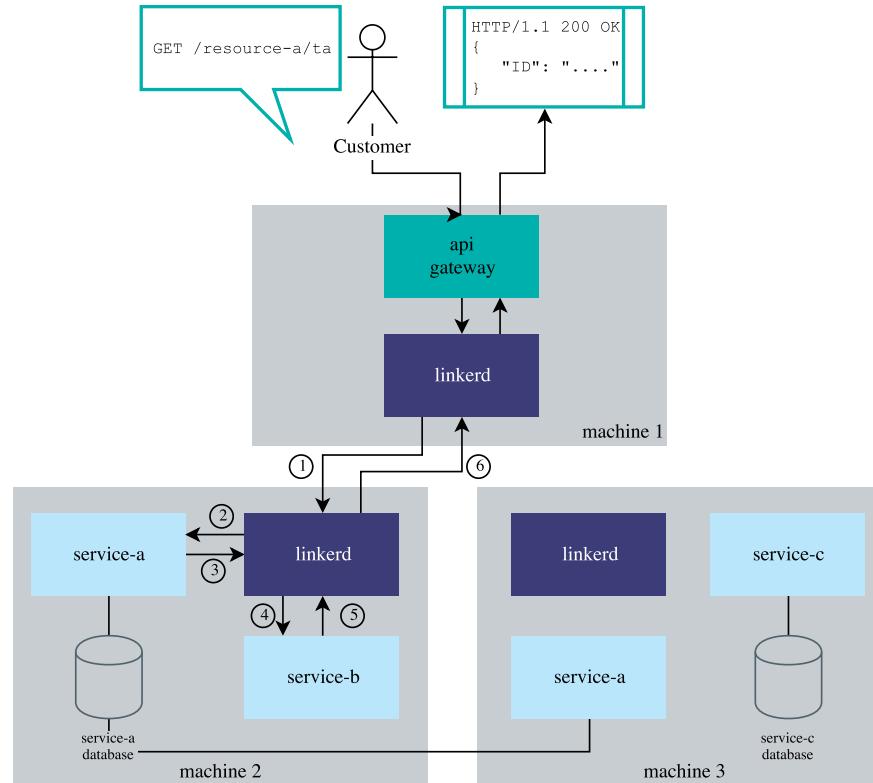
Find a trace

Dependencies

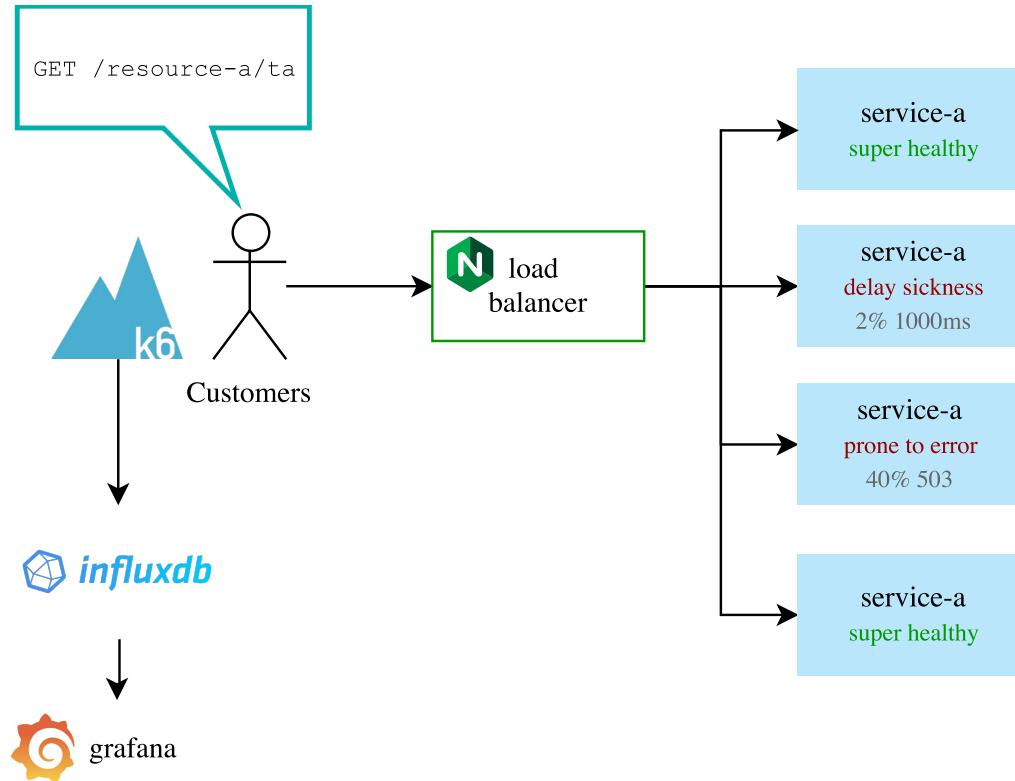
Go to trace



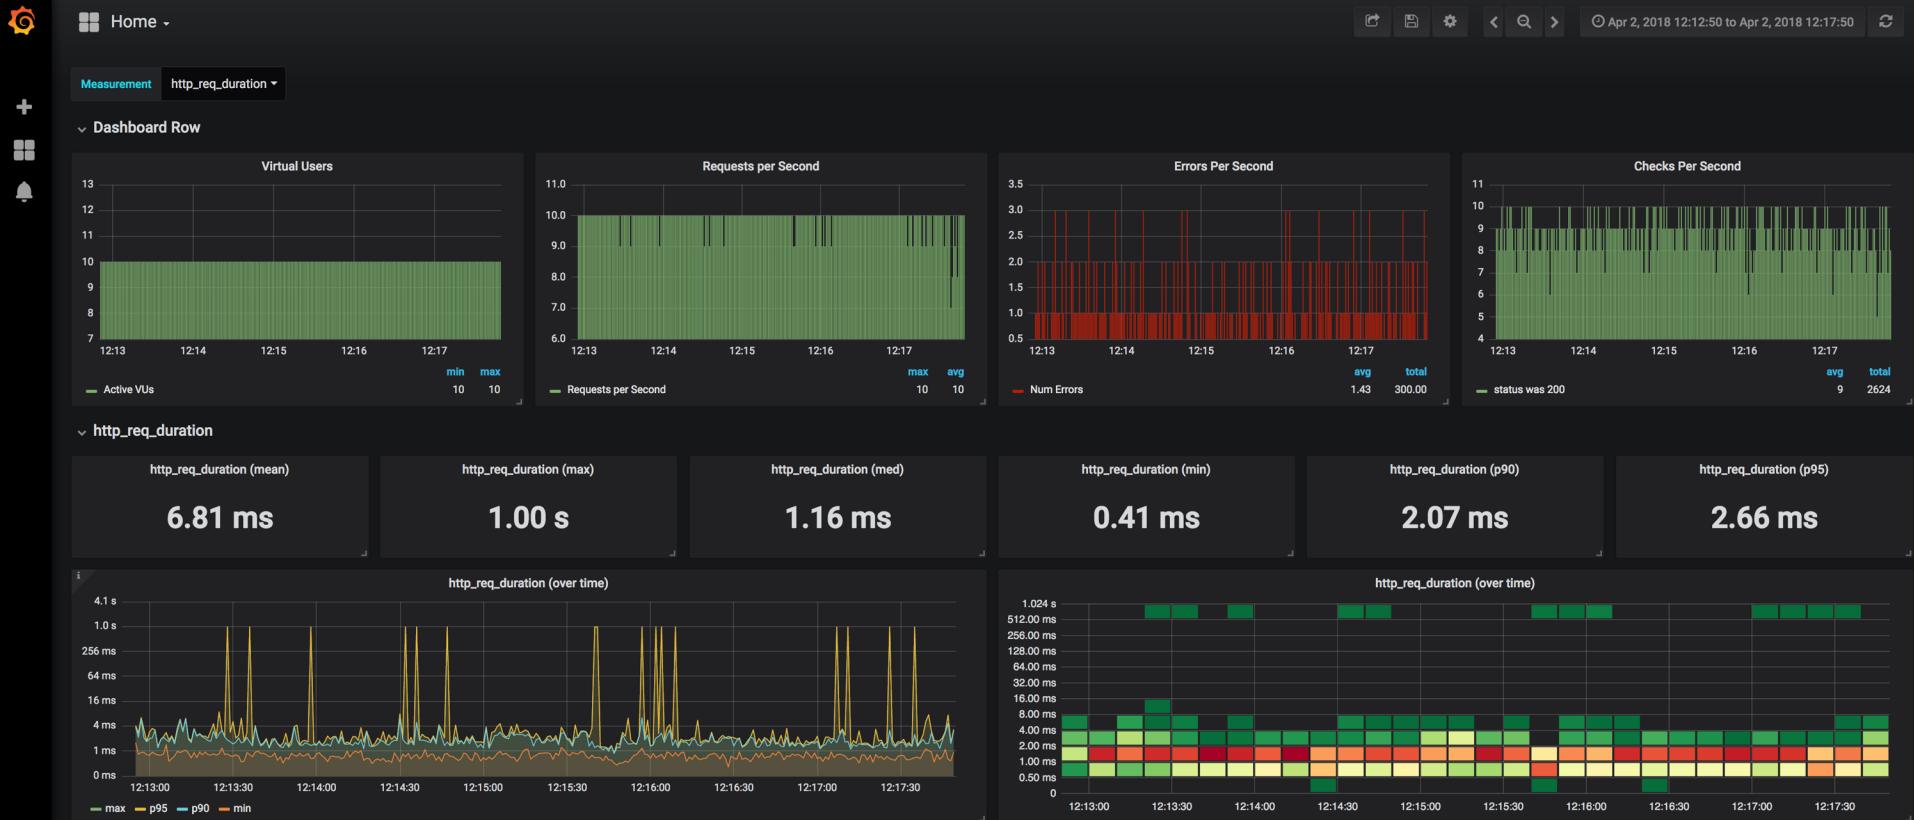
Service mesh in a microservices architecture



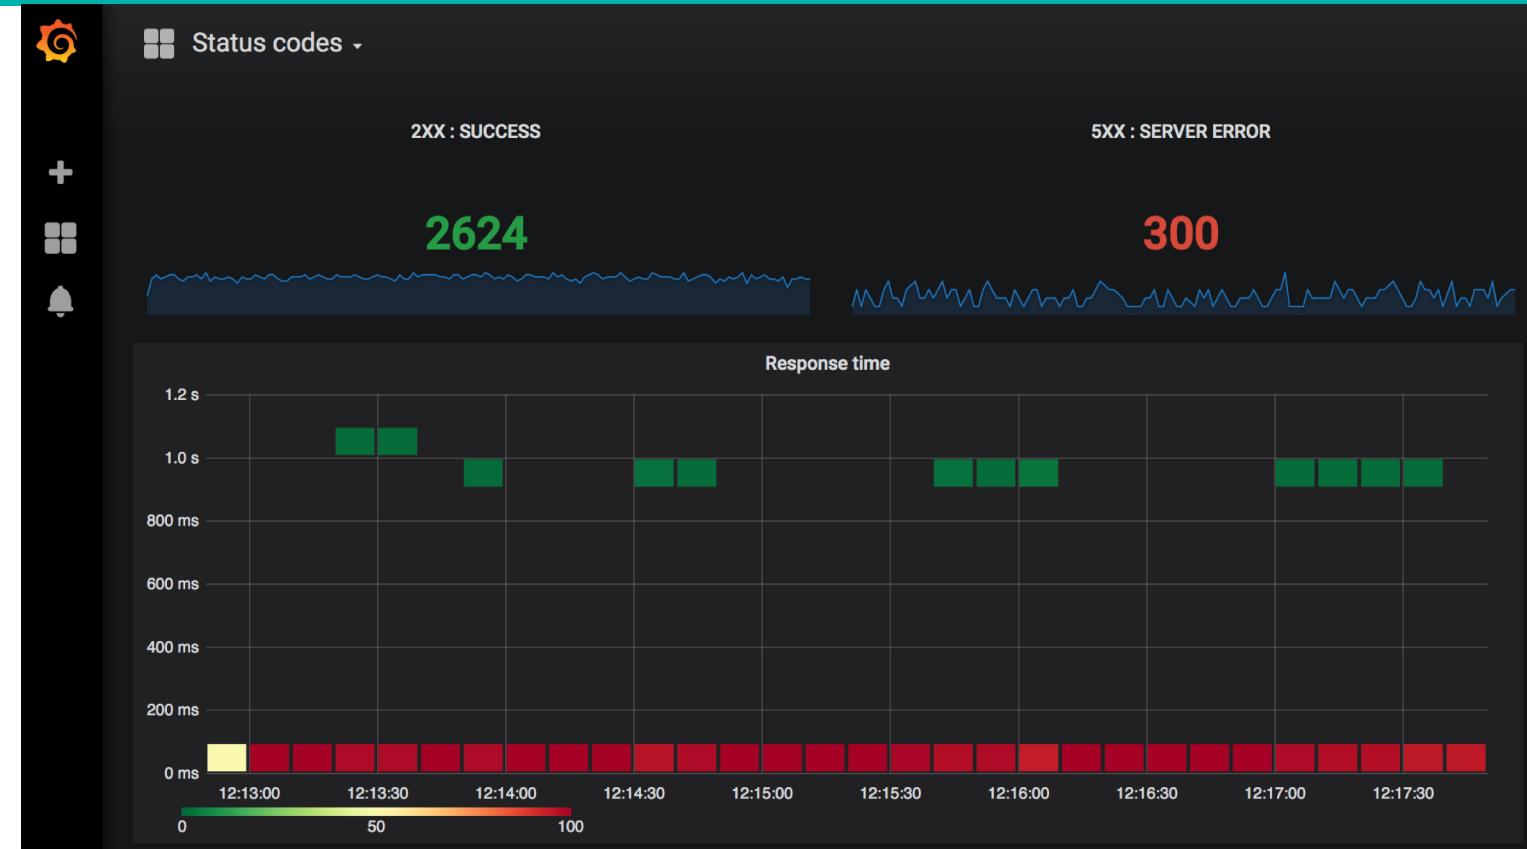
Demo: system-1



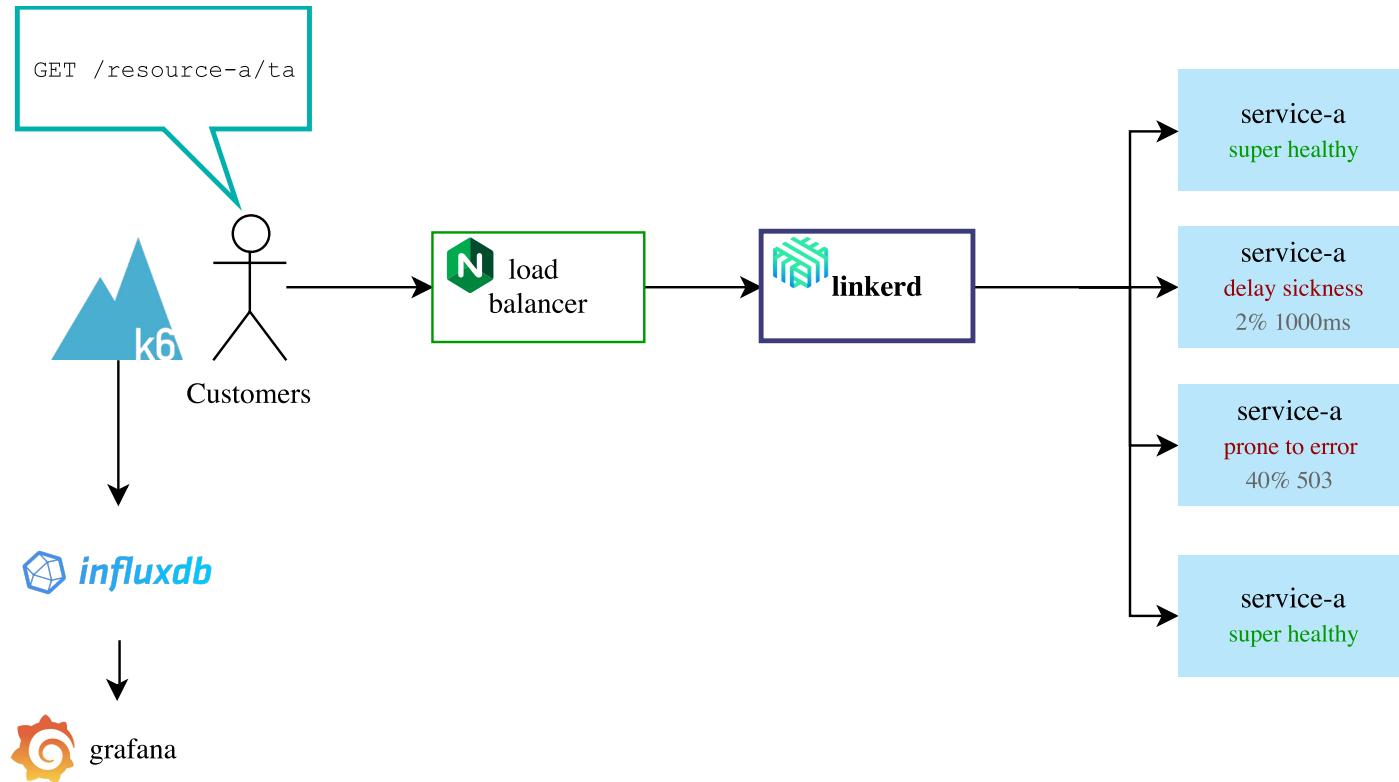
System 1: Response times



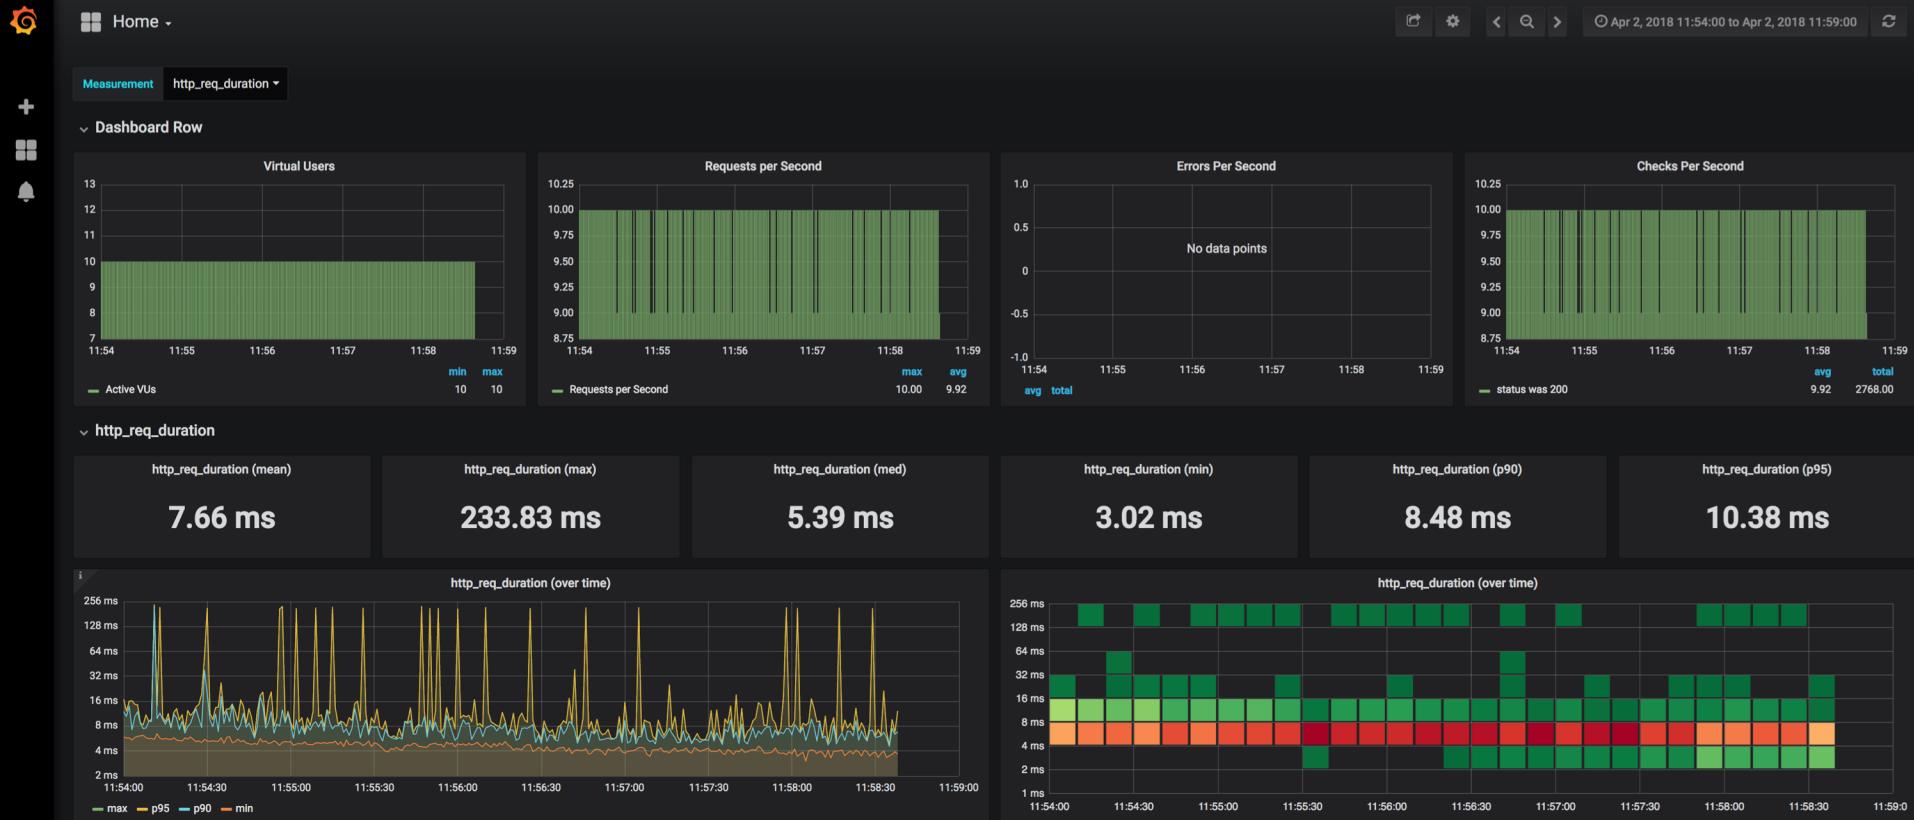
System 1: Response times



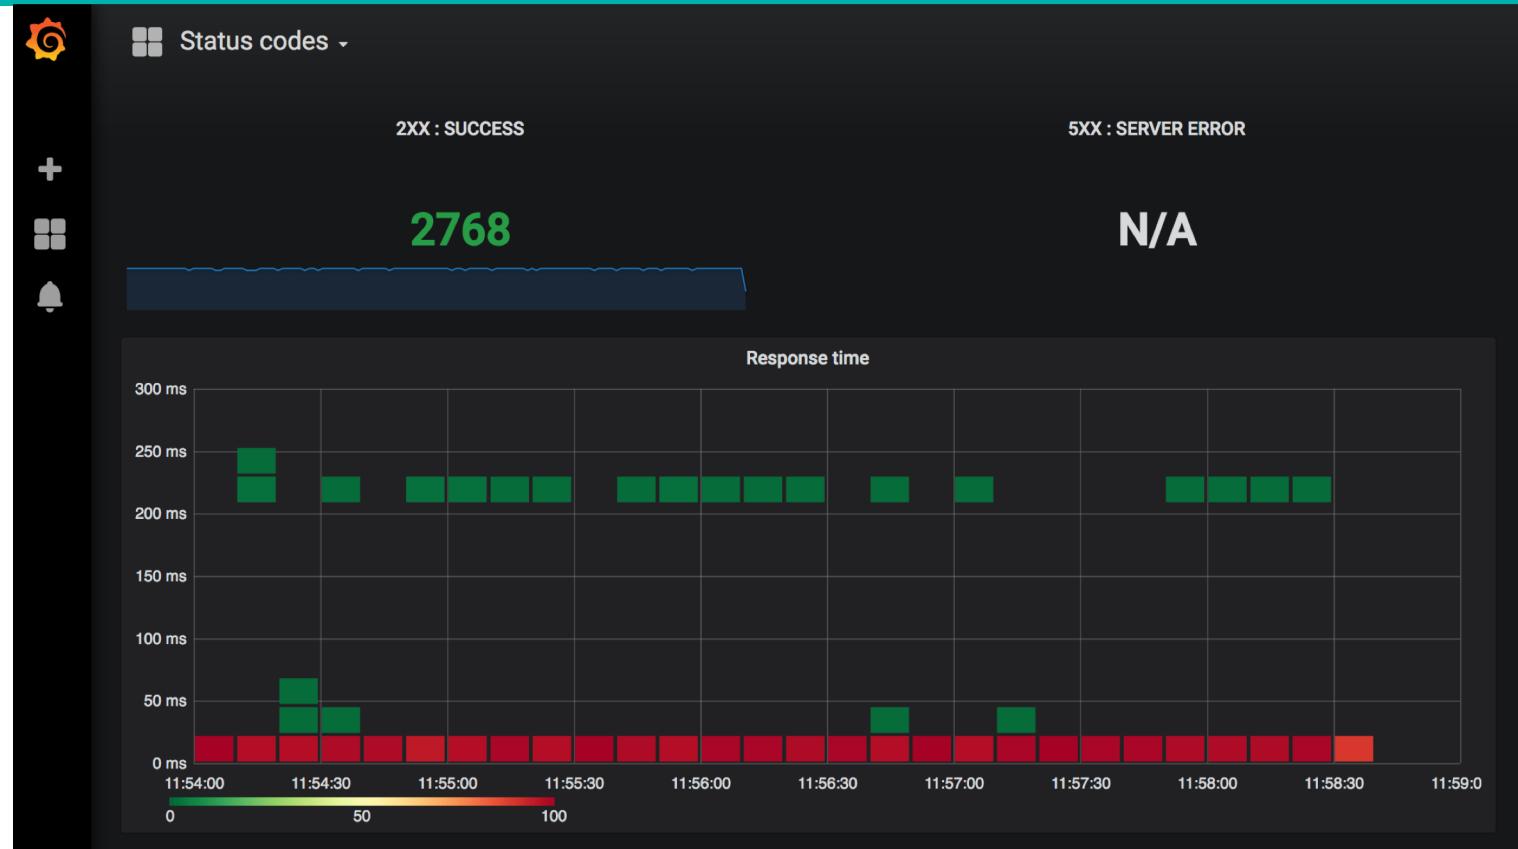
Demo: system-2



System 2: Response times



System 2: Response times



DOI:10.1145/2408776.2408794

Software techniques that tolerate latency variability are vital to building responsive large-scale Web services.

BY JEFFREY DEAN AND LUIZ ANDRÉ BARROSO

The Tail at Scale

as overall use increases. Temporary high-latency episodes (unimportant in moderate-size systems) may come to dominate overall service performance at large scale. Just as fault-tolerant computing aims to create a reliable whole out of less-reliable parts, large online services need to create a predictably responsive whole out of less-predictable parts; we refer to such systems as “latency tail-tolerant,” or simply “tail-tolerant.” Here, we outline some common causes for high-latency episodes in large online services and describe techniques that reduce their severity or mitigate their effect on whole-system performance. In many cases, tail-tolerant techniques can take advantage of resources already deployed to achieve fault-tolerance, resulting in low additional overhead. We explore how these techniques allow system utilization to be driven higher without lengthening the latency tail, thus avoiding wasteful overprovisioning.

Wrap it up

We have used Linkerd to

- Reduce error rates using retries
- Minimize tail latencies using timeouts
- Increase security
- Help diagnose problems using tracing

Thank you

service mesh yay!



<https://github.com/ewilde/kubecon>



this talk

FORM3 <https://form3.tech>

**WE'RE
HIRING!**



and/or



working